

Methodology

- The German Credit dataset was split 70/30 into training and testing sets.
- A logistic regression model was built using 12 key variables.
- To account for the asymmetric cost structure (misclassifying bad customers as good is 5x more costly), we tested multiple probability thresholds (0.30-0.80) and selected the one minimizing total cost: Cost = (False Positives × 5) + (False Negatives × 1).

Solution

Step-1: Load and check out the dataset.

```
#Load data
credit_data <- read.table("germancredit.txt", header = FALSE)
#check data
head(credit_data)
```

```
> head(credit_data)
  V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16 V17 V18 V19 V20 V21
1 A11 6 A34 A43 1169 A65 A75 4 A93 A101 4 A121 67 A143 A152 2 A173 1 A192 A201 1
2 A12 48 A32 A43 5951 A61 A73 2 A92 A101 2 A121 22 A143 A152 1 A173 1 A191 A201 2
3 A14 12 A34 A46 2096 A61 A74 2 A93 A101 3 A121 49 A143 A152 1 A172 2 A191 A201 1
4 A11 42 A32 A42 7882 A61 A74 2 A93 A103 4 A122 45 A143 A153 1 A173 2 A191 A201 1
5 A11 24 A33 A40 4870 A61 A73 3 A93 A101 4 A124 53 A143 A153 2 A173 2 A191 A201 2
6 A14 36 A32 A46 9055 A65 A73 2 A93 A101 4 A124 35 A143 A153 1 A172 2 A192 A201 1
```

Step-2:

The data has 21 columns (V1 to V21). Column V21 tells us if they're good (1) or bad (2) credit. We assume 1=good, 0=bad.

```
#The last column (V21) is our answer: 1 = good credit, 2 = bad credit
#Let's make it 1 = good credit risk, 0 = bad credit risk
credit_data$good_credit <- ifelse(credit_data$V21 == 1, 1, 0)
```

Step-3: Splitting our data into training and testing data.

```
#Split into training and testing data
#We'll use 70% to train, 30% to test
set.seed(123)
n <- nrow(credit_data)
train_indices <- sample(1:n, size = 0.7 * n)

train_data <- credit_data[train_indices, ]
test_data <- credit_data[-train_indices, ]
cat("Training data size:", nrow(train_data), "\n")
cat("Testing data size:", nrow(test_data), "\n\n")
```

Output:

```
> cat("Training data size:", nrow(train_data), "\n")
Training data size: 700
> cat("Testing data size:", nrow(test_data), "\n\n")
Testing data size: 300
```

Step-4: Building logistic regression model with few columns to avoid overfitting issue.

```
#Build the logistic regression model
model <- glm(good_credit ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V9 + V10 + V13 + V14 + V16,
             data = train_data,
             family = binomial(link = "logit"))

#show the model results
cat("== MODEL SUMMARY ==\n")
print(summary(model))
```

Output:

```
== MODEL SUMMARY ==
> print(summary(model))
```

Call:

```
glm(formula = good_credit ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 +
    V9 + V10 + V13 + V14 + V16, family = binomial(link = "logit"),
    data = train_data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.413e-01	9.639e-01	0.147	0.883494
V1A12	2.337e-01	2.561e-01	0.913	0.361468
V1A13	7.405e-01	4.385e-01	1.689	0.091272 .
V1A14	1.645e+00	2.769e-01	5.941	2.83e-09 ***
V2	-3.436e-02	1.043e-02	-3.293	0.000990 ***
V3A31	-3.439e-01	6.625e-01	-0.519	0.603755
V3A32	7.325e-01	5.314e-01	1.379	0.168045
V3A33	7.453e-01	5.904e-01	1.262	0.206834
V3A34	1.674e+00	5.372e-01	3.117	0.001827 **
V4A41	1.647e+00	4.392e-01	3.749	0.000177 ***
V4A410	1.614e+00	9.792e-01	1.649	0.099225 .
V4A42	7.632e-01	3.097e-01	2.464	0.013739 *
V4A43	8.676e-01	2.961e-01	2.930	0.003384 **
V4A44	5.857e-02	9.191e-01	0.064	0.949188
V4A45	-4.920e-01	6.818e-01	-0.722	0.470540
V4A46	2.559e-01	4.889e-01	0.523	0.600638
V4A48	1.807e+00	1.311e+00	1.379	0.168037
V4A49	9.684e-01	3.963e-01	2.444	0.014538 *
V5	-1.156e-04	4.725e-05	-2.446	0.014431 *
V6A62	4.638e-01	3.394e-01	1.366	0.171800
V6A63	8.683e-01	5.071e-01	1.712	0.086856 .
V6A64	1.530e+00	6.200e-01	2.469	0.013567 *
V6A65	1.164e+00	3.278e-01	3.551	0.000384 ***
V8	-3.628e-01	1.029e-01	-3.527	0.000420 ***
V9A92	-1.533e-02	4.458e-01	-0.034	0.972574
V9A93	7.598e-01	4.384e-01	1.733	0.083067 .
V9A94	1.278e-01	5.220e-01	0.245	0.806557
V10A102	-7.055e-01	5.170e-01	-1.365	0.172344
V10A103	1.187e+00	4.986e-01	2.380	0.017305 *
V13	6.654e-03	9.335e-03	0.713	0.476007
V14A142	1.204e-01	5.136e-01	0.234	0.814716
V14A143	2.395e-01	2.902e-01	0.825	0.409304
V16	-2.254e-01	2.402e-01	-0.938	0.348001

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 844.8 on 699 degrees of freedom
 Residual deviance: 630.0 on 667 degrees of freedom
 AIC: 696

Number of Fisher Scoring iterations: 5

Step-5: Make predictions on test data.

```
#Make predictions on TEST data
test_data$predicted_prob <- predict(model, newdata = test_data, type = "response")
```

Step-6: Finding the best threshold, also plotting to check it visually. Bad customers classified as good is 5x worse than good classified as bad.

```
#Find the BEST threshold
# Let's try different thresholds and calculate total cost
thresholds <- seq(0.3, 0.8, by = 0.05)
costs <- numeric(length(thresholds))

for (i in 1:length(thresholds)) {
  threshold <- thresholds[i]
  predicted_class <- ifelse(test_data$predicted_prob >= threshold, 1, 0)

  # False Positive: predicted good (1) but actually bad (0) - COST = 5
  false_positives <- sum(predicted_class == 1 & test_data$good_credit == 0)

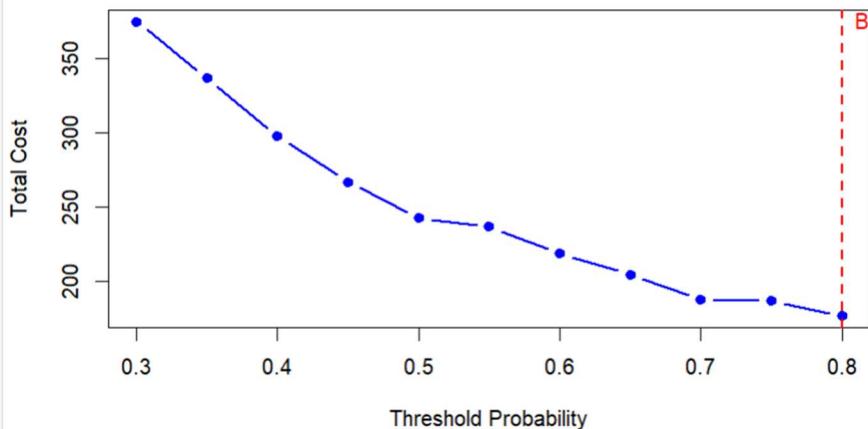
  # False Negative: predicted bad (0) but actually good (1) - COST = 1
  false_negatives <- sum(predicted_class == 0 & test_data$good_credit == 1)

  # Total cost
  costs[i] <- (false_positives * 5) + (false_negatives * 1)
}
```

```
# Find best threshold (minimum cost)
best_index <- which.min(costs)
best_threshold <- thresholds[best_index]

plot(thresholds, costs,
      type = "b",
      col = "blue",
      lwd = 2,
      pch = 19,
      xlab = "Threshold Probability",
      ylab = "Total Cost",
      main = "Finding the Best Threshold")
abline(v = best_threshold, col = "red", lwd = 2, lty = 2)
text(best_threshold, max(costs),
     paste("Best:", best_threshold),
     pos = 4, col = "red")
```

Finding the Best Threshold



We can see that the best threshold is at 0.8 with minimum total cost.

```

cat("\n==== THRESHOLD ANALYSIS ===\n")
result_table <- data.frame(
  Threshold = thresholds,
  Total_Cost = costs
)
print(result_table)

cat("\n*** BEST THRESHOLD:", best_threshold, "***\n")
cat("This minimizes the cost considering bad->good errors are 5x worse\n\n")

#Evaluate model quality with BEST threshold
predicted_class <- ifelse(test_data$predicted_prob >= best_threshold, 1, 0)

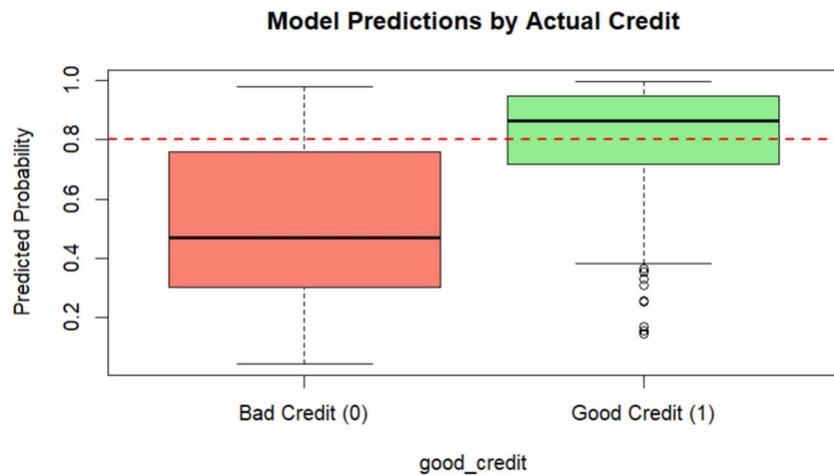
==== THRESHOLD ANALYSIS ====
> result_table <- data.frame(
+   Threshold = thresholds,
+   Total_Cost = costs
+ )
> print(result_table)
  Threshold Total_Cost
1      0.30      375
2      0.35      337
3      0.40      298
4      0.45      267
5      0.50      243
6      0.55      237
7      0.60      219
8      0.65      205
9      0.70      188
10     0.75      187
11     0.80      177
>
> cat("\n*** BEST THRESHOLD:", best_threshold, "***\n")

*** BEST THRESHOLD: 0.8 ***
> cat("This minimizes the cost considering bad->good errors are 5x worse\n\n")
This minimizes the cost considering bad->good errors are 5x worse

```

Step-7: Plotting good vs bad credit predictions using box plot.

```
boxplot(predicted_prob ~ good_credit,
        data = test_data,
        col = c("salmon", "lightgreen"),
        names = c("Bad Credit (0)", "Good Credit (1)"),
        ylab = "Predicted Probability",
        main = "Model Predictions by Actual Credit")
abline(h = best_threshold, col = "red", lwd = 2, lty = 2)
```



Step-8: Creating a confusion matrix to see correct and incorrect predictions.

```
# Confusion Matrix
confusion <- table(Predicted = predicted_class, Actual = test_data$good_credit)
cat("== CONFUSION MATRIX ==\n")
print(confusion)

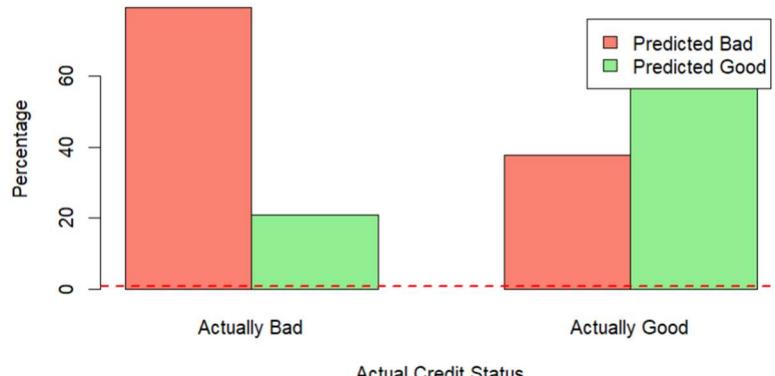
confusion_pct <- prop.table(confusion, margin = 2) * 100
barplot(confusion_pct,
        beside = TRUE,
        col = c("salmon", "lightgreen"),
        legend = c("Predicted Bad", "Predicted Good"),
        xlab = "Actual Credit Status",
        ylab = "Percentage",
        main = "Prediction Accuracy",
        names.arg = c("Actually Bad", "Actually Good"))
```

```
== CONFUSION MATRIX ==
```

```
> print(confusion)
```

Actual		
Predicted	0	1
0	76	77
1	20	127

Prediction Accuracy



This bar chart shows that we missed 20% of bad customers by classifying them as good customers.

Step-9: Calculating metrics

```
# Calculate metrics
accuracy <- sum(predicted_class == test_data$good_credit) / nrow(test_data)
sensitivity <- confusion[2,2] / sum(confusion[,2]) # True positive rate
specificity <- confusion[1,1] / sum(confusion[,1]) # True negative rate

cat("\n==== MODEL QUALITY ===\n")
cat("Accuracy:", round(accuracy * 100, 2), "%\n")
cat("Sensitivity (catching good customers):", round(sensitivity * 100, 2), "%\n")
cat("Specificity (catching bad customers):", round(specificity * 100, 2), "%\n")

#show what the coefficients mean
cat("\n==== WHAT THE MODEL LEARNED (Simplified) ===\n")
cat("Positive coefficient = increases chance of being good credit\n")
cat("Negative coefficient = decreases chance of being good credit\n")
cat("Bigger number = stronger effect\n\n")

coefs <- summary(model)$coefficients[,1]
cat("Top factors:\n")
print(round(coefs, 3))

==== MODEL QUALITY ====
> cat("Accuracy:", round(accuracy * 100, 2), "%\n")
Accuracy: 67.67 %
> cat("Sensitivity (catching good customers):", round(sensitivity * 100, 2), "%\n")
Sensitivity (catching good customers): 62.25 %
> cat("Specificity (catching bad customers):", round(specificity * 100, 2), "%\n")
Specificity (catching bad customers): 79.17 %
>
> #show what the coefficients mean
> cat("\n==== WHAT THE MODEL LEARNED (Simplified) ===\n")

==== WHAT THE MODEL LEARNED (Simplified) ====
> cat("Positive coefficient = increases chance of being good credit\n")
Positive coefficient = increases chance of being good credit
> cat("Negative coefficient = decreases chance of being good credit\n")
Negative coefficient = decreases chance of being good credit
> cat("Bigger number = stronger effect\n\n")
Bigger number = stronger effect

>
> coefs <- summary(model)$coefficients[,1]
> cat("Top factors:\n")
Top factors:
> print(round(coefs, 3))
(Intercept) V1A12 V1A13 V1A14 V2 V3A31 V3A32
0.141 0.234 0.741 1.645 -0.034 -0.344 0.733
V3A33 V3A34 V4A41 V4A410 V4A42 V4A43 V4A44
0.745 1.674 1.647 1.614 0.763 0.868 0.059
V4A45 V4A46 V4A48 V4A49 V5 V6A62 V6A63
-0.492 0.256 1.807 0.968 0.000 0.464 0.868
V6A64 V6A65 V8 V9A92 V9A93 V9A94 V10A102
1.530 1.164 -0.363 -0.015 0.760 0.128 -0.706
V10A103 V13 V14A142 V14A143 V16
1.187 0.007 0.120 0.239 -0.225
```

1. Accuracy: `accuracy <- sum(predicted_class == test_data$good_credit) / nrow(test_data)`

67.67% accuracy = we're correct approximately 68 times out of 100

Sensitivity + Specificity show the full picture of what types of mistakes we're making

2. Sensitivity: `sensitivity <- confusion[2,2] / sum(confusion[,2])`

How good we are at finding GOOD customers?

Sensitivity = 62.25 %

3. Specificity: `specificity <- confusion[1,1] / sum(confusion[,1])`

How good we are at catching BAD customers?

Specificity = 79.17 %

It is higher because we'd rather reject a good person than approve a bad one (due to the 5x cost)

Result Discussion:

The logistic regression model achieved 67.67% accuracy with specificity (79.17%) notably higher than sensitivity (62.25%), reflecting the cost-optimized threshold that prioritizes catching bad customers over approving all good ones. This conservative approach aligns with the 5:1 cost ratio, effectively minimizing expensive false approvals (bad customers classified as good) at the cost of rejecting some creditworthy applicants.